

# 1. Read the data

```
-----
FileNotFoundError                                Traceback (most recent call last)
Cell In[10], line 1
----> 1 df = pd.read_csv(r'notebook\data\depression_data.csv')

File c:\ML\mentalhealthapp\venv\lib\site-packages\pandas\io\parsers\readers.py:912, in read_csv(filepath_or_buffer, sep, delimiter, header, names, index_col, usecols, dtype, engine, converters, true_values, false_values, skipinitialspace, skiprows, skipfooter, nrows, na_values, keep_default_na, na_filter, verbose, skip_blank_lines, parse_dates, infer_datetime_format, keep_date_col, date_parser, date_format, dayfirst, cache_dates, iterator, chunksize, compression, thousands, decimal, lineterminator, quotechar, quoting, doublequote, escapechar, comment, encoding, encoding_errors, dialect, on_bad_lines, delim_whitespace, low_memory, memory_map, float_precision, storage_options, dtype_backend)
    899 kwds_defaults = _refine_defaults_read(
    900     dialect,
    901     delimiter,
    (...)
    908     dtype_backend=dtype_backend,
    909 )
    910 kwds.update(kwds_defaults)
--> 912 return _read(filepath_or_buffer, kwds)

File c:\ML\mentalhealthapp\venv\lib\site-packages\pandas\io\parsers\readers.py:577, in _read(filepath_or_buffer, kwds)
    574 _validate_names(kwds.get("names", None))
    576 # Create the parser.
--> 577 parser = TextFileReader(filepath_or_buffer, **kwds)
    579 if chunksize or iterator:
    580     return parser

File c:\ML\mentalhealthapp\venv\lib\site-packages\pandas\io\parsers\readers.py:1407, in TextFileReader.__init__(self, f, engine, **kwds)
    1404 self.options["has_index_names"] = kwds["has_index_names"]
    1406 self.handles: IOHandles | None = None
-> 1407 self._engine = self._make_engine(f, self.engine)

File c:\ML\mentalhealthapp\venv\lib\site-packages\pandas\io\parsers\readers.py:1661, in TextFileReader._make_engine(self, f, engine)
    1659 if "b" not in mode:
    1660     mode += "b"
-> 1661 self.handles = get_handle(
    1662     f,
    1663     mode,
    1664     encoding=self.options.get("encoding", None),
    1665     compression=self.options.get("compression", None),
    1666     memory_map=self.options.get("memory_map", False),
    1667     is_text=is_text,
    1668     errors=self.options.get("encoding_errors", "strict"),
    1669     storage_options=self.options.get("storage_options", None),
    1670 )
    1671 assert self.handles is not None
    1672 f = self.handles.handle

File c:\ML\mentalhealthapp\venv\lib\site-packages\pandas\io\common.py:859, in get_handle(path_or_buffer, mode, encoding, compression, memory_map, is_text, errors, storage_options)
    854 elif isinstance(handle, str):
    855     # Check whether the filename is to be opened in binary mode.
```

```
856 # Binary mode does not support 'encoding' and 'newline'.
857 if ioargs.encoding and "b" not in ioargs.mode:
858     # Encoding
--> 859     handle = open(
860         handle,
861         ioargs.mode,
862         encoding=ioargs.encoding,
863         errors=errors,
864         newline="",
865     )
866 else:
867     # Binary mode
868     handle = open(handle, ioargs.mode)

FileNotFoundError: [Errno 2] No such file or directory: 'notebook\\data\\depression_data.csv'
```

## 2. Get the look & feel of the data

|   | Name             | Age | Marital Status | Education Level   | Number of Children | Smoking Status | Physical Activity Level | Employment Status | Income    | Alcohol Consumption | Dietary Habits | Pat |
|---|------------------|-----|----------------|-------------------|--------------------|----------------|-------------------------|-------------------|-----------|---------------------|----------------|-----|
| 0 | Christine Barker | 31  | Married        | Bachelor's Degree | 2                  | Non-smoker     | Active                  | Unemployed        | 26265.67  | Moderate            | Moderate       |     |
| 1 | Jacqueline Lewis | 55  | Married        | High School       | 1                  | Non-smoker     | Sedentary               | Employed          | 42710.36  | High                | Unhealthy      |     |
| 2 | Shannon Church   | 78  | Widowed        | Master's Degree   | 1                  | Non-smoker     | Sedentary               | Employed          | 125332.79 | Low                 | Unhealthy      | (   |
| 3 | Charles Jordan   | 58  | Divorced       | Master's Degree   | 3                  | Non-smoker     | Moderate                | Unemployed        | 9992.78   | Moderate            | Moderate       |     |
| 4 | Michael Rich     | 18  | Single         | High School       | 0                  | Non-smoker     | Sedentary               | Unemployed        | 8595.08   | Low                 | Moderate       |     |

|        | Name            | Age | Marital Status | Education Level   | Number of Children | Smoking Status | Physical Activity Level | Employment Status | Income    | Alcohol Consumption | Dietary Habits |  |
|--------|-----------------|-----|----------------|-------------------|--------------------|----------------|-------------------------|-------------------|-----------|---------------------|----------------|--|
| 413763 | Sean Miller     | 68  | Married        | Master's Degree   | 0                  | Former         | Moderate                | Employed          | 109233.43 | Low                 | Healthy        |  |
| 413764 | Christina Brown | 26  | Single         | Bachelor's Degree | 0                  | Current        | Active                  | Employed          | 96760.97  | Low                 | Healthy        |  |
| 413765 | Matthew Jenkins | 57  | Married        | Bachelor's Degree | 0                  | Non-smoker     | Sedentary               | Employed          | 77353.26  | Moderate            | Moderate       |  |
| 413766 | Gary Faulkner   | 71  | Married        | Associate Degree  | 2                  | Non-smoker     | Sedentary               | Unemployed        | 24557.08  | Moderate            | Moderate       |  |
| 413767 | Joseph Johnson  | 62  | Widowed        | Master's Degree   | 0                  | Former         | Moderate                | Employed          | 107125.74 | Moderate            | Healthy        |  |

## 3. Understand the datatypes & basic statistics of the data

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 413768 entries, 0 to 413767
Data columns (total 16 columns):
#      Column                                Non-Null Count  Dtype
---  -
0      Name                                413768 non-null  object
1      Age                                413768 non-null  int64
2      Marital Status                      413768 non-null  object
3      Education Level                    413768 non-null  object
4      Number of Children                 413768 non-null  int64
5      Smoking Status                     413768 non-null  object
6      Physical Activity Level            413768 non-null  object
7      Employment Status                 413768 non-null  object
8      Income                             413768 non-null  float64
9      Alcohol Consumption                413768 non-null  object
10     Dietary Habits                     413768 non-null  object
11     Sleep Patterns                     413768 non-null  object
12     History of Mental Illness          413768 non-null  object
13     History of Substance Abuse         413768 non-null  object
14     Family History of Depression       413768 non-null  object
15     Chronic Medical Conditions         413768 non-null  object
dtypes: float64(1), int64(2), object(13)
memory usage: 50.5+ MB
```

|       | Age           | Number of Children | Income        |
|-------|---------------|--------------------|---------------|
| count | 413768.000000 | 413768.000000      | 413768.000000 |
| mean  | 49.000713     | 1.298972           | 50661.707971  |
| std   | 18.158759     | 1.237054           | 40624.100565  |
| min   | 18.000000     | 0.000000           | 0.410000      |
| 25%   | 33.000000     | 0.000000           | 21001.030000  |
| 50%   | 49.000000     | 1.000000           | 37520.135000  |
| 75%   | 65.000000     | 2.000000           | 76616.300000  |
| max   | 80.000000     | 4.000000           | 209995.220000 |

```
No      287943
Yes     125825
Name: History of Mental Illness, dtype: int64

Marital Status      False
Education Level     False
Number of Children  False
Smoking Status      False
Physical Activity Level False
Employment Status   False
Income              False
Alcohol Consumption False
Dietary Habits      False
Sleep Patterns      False
History of Mental Illness False
History of Substance Abuse False
Family History of Depression False
Chronic Medical Conditions False
Gender              False
income_groups       False
age_groups          False
dtype: bool
```

False

# Feature Engineering.


## Generate gender from Name

```
male      209232
female    204536
Name: Gender, dtype: int64
```

209995.22


| History of Mental Illness | No     | Yes   |
|---------------------------|--------|-------|
| Marital Status            |        |       |
| Divorced                  | 23111  | 9618  |
| Married                   | 167155 | 73289 |
| Single                    | 50690  | 21420 |
| Widowed                   | 46987  | 21498 |

```
(array([103438., 125739., 56380., 39450., 43574., 18596., 14751.,
        7515., 2179., 2146.]),
 array([4.10000000e-01, 2.09998910e+04, 4.19993720e+04, 6.29988530e+04,
        8.39983340e+04, 1.04997815e+05, 1.25997296e+05, 1.46996777e+05,
        1.67996258e+05, 1.88995739e+05, 2.09995220e+05])),
<BarContainer object of 10 artists>)
```





```
C:\Users\admin\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot`
` is a deprecated function and will be removed in a future version. Please adapt your code to use
either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level f
unction for histograms).
warnings.warn(msg, FutureWarning)
```

```
<AxesSubplot:xlabel='Income', ylabel='Density'>
```



1.086114667512213



The Income feature has the outliers & needs to be treated. I would like to go with capping method.

21001.03  
76616.3  
55615.270000000004  
Upper limit 160039.20500000002  
Lower limit -62421.875

|        | Age | Marital Status | Education Level | Number of Children | Smoking Status | Physical Activity Level | Employment Status | Income    | Alcohol Consumption | Dietary Habits | Sleep Patterns |
|--------|-----|----------------|-----------------|--------------------|----------------|-------------------------|-------------------|-----------|---------------------|----------------|----------------|
| 15     | 38  | Married        | PhD             | 0                  | Non-smoker     | Moderate                | Employed          | 202449.17 | High                | Healthy        | Fair           |
| 105    | 53  | Widowed        | PhD             | 3                  | Former         | Sedentary               | Employed          | 169400.38 | High                | Unhealthy      | Good           |
| 170    | 36  | Divorced       | PhD             | 3                  | Non-smoker     | Sedentary               | Employed          | 180084.56 | High                | Unhealthy      | Fair           |
| 193    | 64  | Married        | PhD             | 1                  | Non-smoker     | Moderate                | Employed          | 193843.44 | Low                 | Healthy        | Poor           |
| 319    | 31  | Married        | PhD             | 1                  | Current        | Moderate                | Employed          | 177029.40 | High                | Unhealthy      | Poor           |
| ...    | ... | ...            | ...             | ...                | ...            | ...                     | ...               | ...       | ...                 | ...            | ...            |
| 413461 | 77  | Widowed        | PhD             | 0                  | Non-smoker     | Moderate                | Employed          | 200828.61 | Moderate            | Moderate       | Fair           |
| 413563 | 76  | Married        | PhD             | 1                  | Former         | Sedentary               | Employed          | 164436.49 | High                | Unhealthy      | Poor           |
| 413574 | 41  | Divorced       | PhD             | 0                  | Former         | Sedentary               | Employed          | 171921.83 | Low                 | Moderate       | Good           |
| 413577 | 30  | Married        | PhD             | 1                  | Non-smoker     | Moderate                | Employed          | 169051.46 | High                | Healthy        | Poor           |
| 413754 | 34  | Married        | PhD             | 4                  | Non-smoker     | Active                  | Employed          | 185657.43 | Moderate            | Moderate       | Fair           |

5157 rows × 19 columns

| Age | Marital Status | Education Level | Number of Children | Smoking Status | Physical Activity Level | Employment Status | Income | Alcohol Consumption | Dietary Habits | Sleep Patterns | History of Mental Illness | Hist Sub: |
|-----|----------------|-----------------|--------------------|----------------|-------------------------|-------------------|--------|---------------------|----------------|----------------|---------------------------|-----------|
|-----|----------------|-----------------|--------------------|----------------|-------------------------|-------------------|--------|---------------------|----------------|----------------|---------------------------|-----------|

array(['PhD'], dtype=object)

All the outliers in Income column are PHD holders & rightly so as they have the highest educational degree & are generally paid higher compared to other degree holders.

```
C:\Users\admin\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot`
` is a deprecated function and will be removed in a future version. Please adapt your code to use
either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level f
unction for histograms).
    warnings.warn(msg, FutureWarning)
C:\Users\admin\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the foll
owing variable as a keyword arg: x. From version 0.12, the only valid positional argument will be
`data`, and passing other arguments without an explicit keyword will result in an error or misinte
rpretation.
    warnings.warn(
```




C:\Users\admin\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\admin\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```


```
<AxesSubplot:xlabel='Income'>
```



C:\Users\admin\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


```
warnings.warn(
```

```
<AxesSubplot:xlabel='Age'>
```

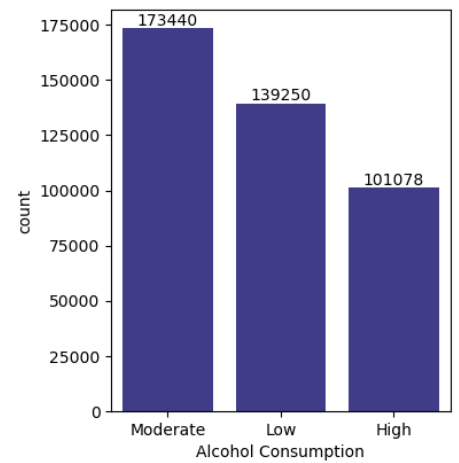

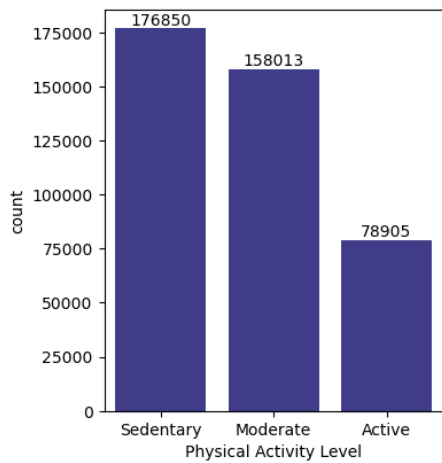
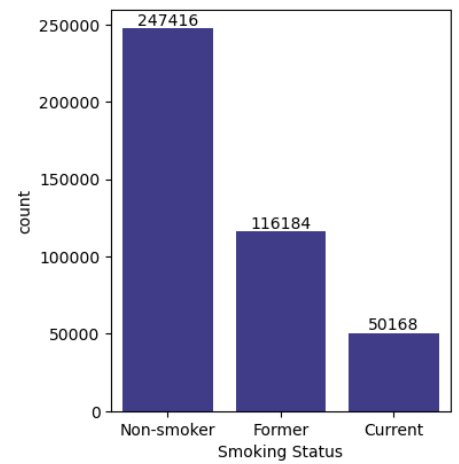

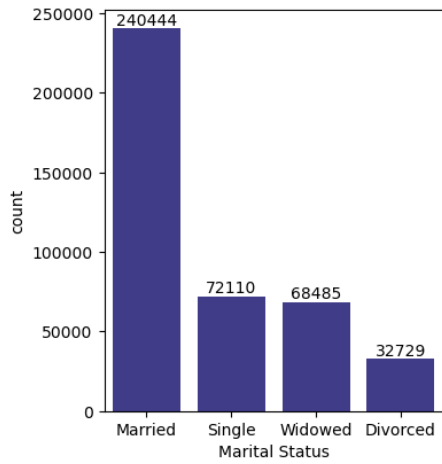


```
array(['Moderate', 'Unhealthy', 'Healthy'], dtype=object)
```

# Univariate & Bi-variate Analysis




The gender feature seems to be having equal distribution. We can check the correlation between this feature with y variable .







Elderly 164709  
Retired 150871  
Adults 59123  
Young Adults 39065  
Name: age\_groups, dtype: int64




Elderly people have higher risk of mental illness compared to young people.



Financially weaker group of people have higher risk of mental illness compared to rich people.

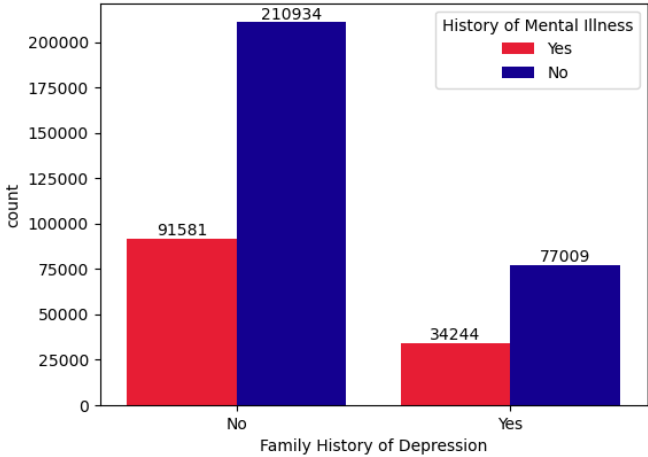
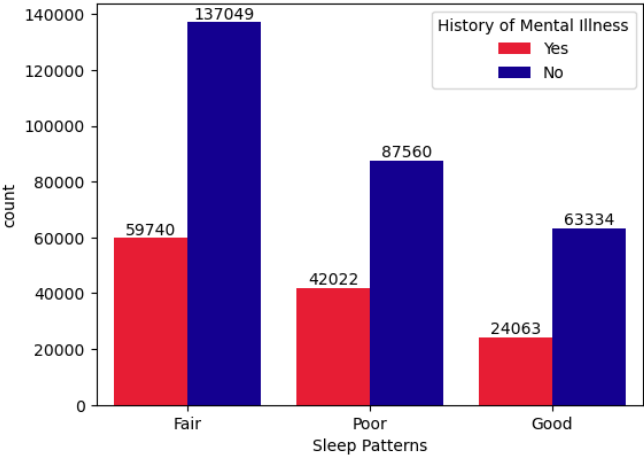




It seems like both males, females have similar risk of health issues.

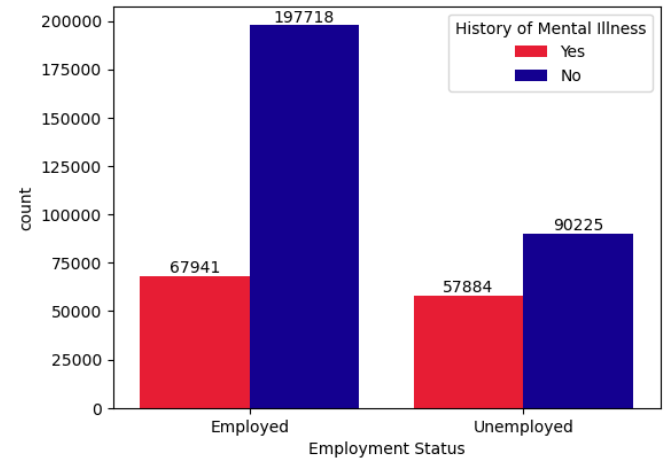
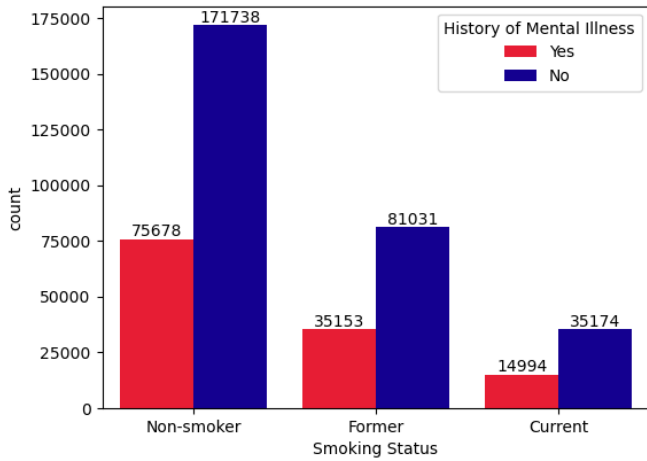
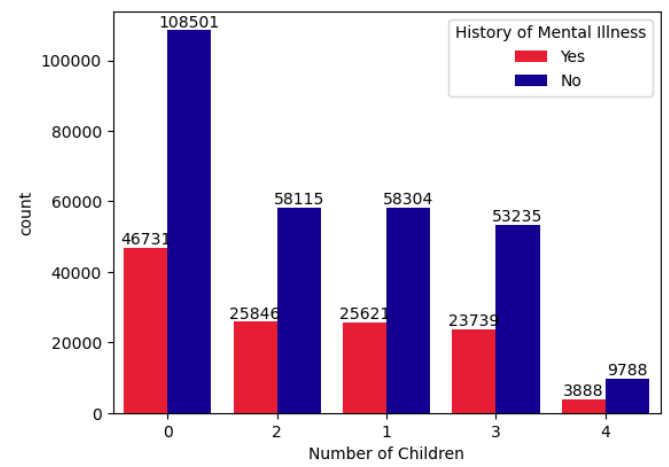



People with higher degree like PhD, Master's have lower risk of mental illness relative to people with High school or

Bachelor's degree.


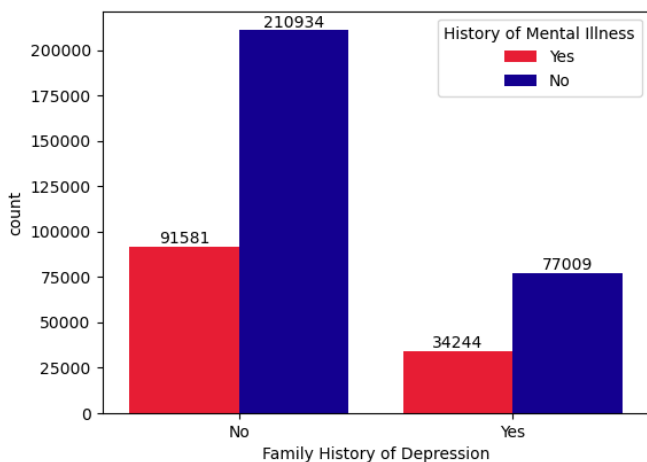
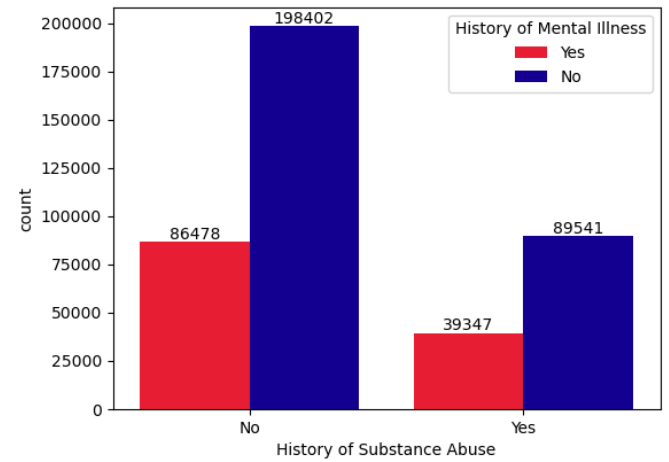
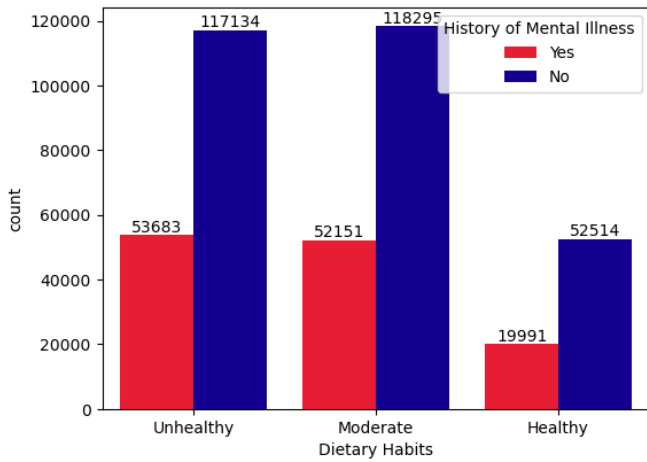


People with Active lifestyle, Good sleeping pattern, High alcohol consumption have lower risk of mental illness.



People with no children , Married people , unemployed have higher chance of mental illness.

Divorced people tend to have less risk as well as the current-moker have less risk.



People who follow healthy lifestyle tend have less risk of mental illness.

# Correlation between the categorical variables using chi-square test

Correlated Categorical Features:

Marital Status and History of Mental Illness are correlated with p-value 5.409031887961224e-14

Education Level and History of Mental Illness are correlated with p-value 5.678973600047682e-302

Smoking Status and History of Mental Illness are correlated with p-value 0.0032730141384533966

Physical Activity Level and History of Mental Illness are correlated with p-value 2.5562562939064635e-07

Employment Status and History of Mental Illness are correlated with p-value 0.0

Alcohol Consumption and History of Mental Illness are correlated with p-value 2.7124703085166094e-17

Dietary Habits and History of Mental Illness are correlated with p-value 2.1242586440978167e-79

Sleep Patterns and History of Mental Illness are correlated with p-value 3.30469070431477e-129

History of Mental Illness and Family History of Depression are correlated with p-value 0.0016897082070443303

History of Mental Illness and Chronic Medical Conditions are correlated with p-value 0.00010435052869828637

History of Mental Illness and income\_groups are correlated with p-value 0.0

History of Mental Illness and age\_groups are correlated with p-value 1.4262358344020623e-64

Independent Categorical Features:

History of Mental Illness and History of Substance Abuse are independent with p-value 0.2666100524904718

History of Mental Illness and Fname are independent with p-value 0.11847747615352716

History of Mental Illness and Gender are independent with p-value 0.4571005011796979

History of Mental Illness is not correlated with Name , Gender & History of Substance Abuse.

We can filter these features.

```
No      287943
Yes      125825
Name: History of Mental Illness, dtype: int64

array([ 0.          ,  1.          ,  0.          ,  0.          ,  0.          ,
        1.          ,  0.          ,  0.          ,  0.          ,  0.          ,
        0.          ,  1.          ,  1.          ,  0.          ,  0.          ,
        0.          ,  1.          ,  0.          ,  0.          ,  1.          ,
        0.          ,  1.          ,  0.          ,  1.          ,  0.          ,
        0.          ,  1.          ,  0.          ,  0.          ,  1.          ,
        0.          ,  0.          ,  0.          ,  0.          ,  1.          ,
        0.          ,  0.          ,  0.          ,  0.5666923 , -0.60745955])

array([1, 1, 0, ..., 0, 0, 0])

125825
```

# Trying Classifier models on un-balanced raw data.

Training Logistic Regression...

```
C:\Users\admin\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\admin\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\admin\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\admin\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

Logistic Regression Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.70      | 1.00   | 0.82     | 86319   |
| 1            | 0.00      | 0.00   | 0.00     | 37812   |
| accuracy     |           |        | 0.70     | 124131  |
| macro avg    | 0.35      | 0.50   | 0.41     | 124131  |
| weighted avg | 0.48      | 0.70   | 0.57     | 124131  |

Logistic Regression Confusion Matrix:

```
[[86319  0]
 [37812  0]]
```

-----

Training Decision Tree...

Decision Tree Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.70      | 0.69   | 0.70     | 86319   |
| 1            | 0.32      | 0.34   | 0.33     | 37812   |
| accuracy     |           |        | 0.58     | 124131  |
| macro avg    | 0.51      | 0.51   | 0.51     | 124131  |
| weighted avg | 0.59      | 0.58   | 0.58     | 124131  |

Decision Tree Confusion Matrix:

```
[[59236 27083]
 [24879 12933]]
```

-----

Training Random Forest...

Random Forest Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.70      | 0.83   | 0.76     | 86319   |
| 1            | 0.34      | 0.19   | 0.25     | 37812   |
| accuracy     |           |        | 0.64     | 124131  |
| macro avg    | 0.52      | 0.51   | 0.50     | 124131  |
| weighted avg | 0.59      | 0.64   | 0.61     | 124131  |

Random Forest Confusion Matrix:

```
[[72028 14291]
 [30529  7283]]
```

-----  
Training Gradient Boosting (XGBoost)...

Gradient Boosting (XGBoost) Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.70      | 0.99   | 0.82     | 86319   |
| 1            | 0.39      | 0.01   | 0.02     | 37812   |
| accuracy     |           |        | 0.69     | 124131  |
| macro avg    | 0.54      | 0.50   | 0.42     | 124131  |
| weighted avg | 0.60      | 0.69   | 0.58     | 124131  |

Gradient Boosting (XGBoost) Confusion Matrix:

```
[[85707  612]
 [37421  391]]
```

-----  
Training Naive Bayes...

Naive Bayes Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.75      | 0.62   | 0.68     | 86319   |
| 1            | 0.37      | 0.51   | 0.43     | 37812   |
| accuracy     |           |        | 0.59     | 124131  |
| macro avg    | 0.56      | 0.57   | 0.56     | 124131  |
| weighted avg | 0.63      | 0.59   | 0.60     | 124131  |

Naive Bayes Confusion Matrix:

```
[[53746 32573]
 [18341 19471]]
```

-----  
Training LightGBM...

```
[LightGBM] [Info] Number of positive: 88013, number of negative: 201624
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.043929 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 390
[LightGBM] [Info] Number of data points in the train set: 289637, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303873 -> initscore=-0.828920
[LightGBM] [Info] Start training from score -0.828920
```

LightGBM Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.70      | 1.00   | 0.82     | 86319   |
| 1            | 0.42      | 0.00   | 0.00     | 37812   |
| accuracy     |           |        | 0.70     | 124131  |
| macro avg    | 0.56      | 0.50   | 0.41     | 124131  |
| weighted avg | 0.61      | 0.70   | 0.57     | 124131  |

LightGBM Confusion Matrix:

```
[[86308  11]
 [37804  8]]
```

|         |                             |          |          |          |          |
|---------|-----------------------------|----------|----------|----------|----------|
| 0       | Logistic Regression         | 0.695386 | 0.483562 | 0.695386 | 0.570445 |
| 1       | Decision Tree               | 0.581394 | 0.588159 | 0.581394 | 0.584614 |
| 2       | Random Forest               | 0.638930 | 0.591217 | 0.638930 | 0.605087 |
| 3       | Gradient Boosting (XGBoost) | 0.693606 | 0.602793 | 0.693606 | 0.575250 |
| 4       | Naive Bayes                 | 0.589837 | 0.632424 | 0.589837 | 0.603893 |
| 5       | LightGBM                    | 0.695362 | 0.611833 | 0.695362 | 0.570552 |
| AUC-ROC |                             |          |          |          |          |
| 0       |                             | 0.593136 |          |          |          |
| 1       |                             | 0.513946 |          |          |          |
| 2       |                             | 0.544688 |          |          |          |
| 3       |                             | 0.593340 |          |          |          |
| 4       |                             | 0.591918 |          |          |          |
| 5       |                             | 0.598066 |          |          |          |

## Balancing data using SMOTE

287943

287943

Training Logistic Regression...

|  |           |        |          |         |  |
|--|-----------|--------|----------|---------|--|
| Logistic Regression Classification Report: |           |        |          |         |  |
|  | precision | recall | f1-score | support |  |
| 0  | 0.57      | 0.60   | 0.58     | 86344   |  |
| 1  | 0.58      | 0.55   | 0.56     | 86422   |  |
| accuracy                                   |           |        | 0.57     | 172766  |  |
| macro avg                                  | 0.57      | 0.57   | 0.57     | 172766  |  |
| weighted avg                               | 0.57      | 0.57   | 0.57     | 172766  |  |

Logistic Regression Confusion Matrix:  
[[51439 34905]  
[39058 47364]]

-----  
Training Decision Tree...

|                                      |           |        |          |         |  |
|--------------------------------------|-----------|--------|----------|---------|--|
| Decision Tree Classification Report: |           |        |          |         |  |
|                                      | precision | recall | f1-score | support |  |
| 0                                    | 0.65      | 0.64   | 0.64     | 86344   |  |
| 1                                    | 0.65      | 0.66   | 0.65     | 86422   |  |
| accuracy                             |           |        | 0.65     | 172766  |  |
| macro avg                            | 0.65      | 0.65   | 0.65     | 172766  |  |
| weighted avg                         | 0.65      | 0.65   | 0.65     | 172766  |  |

Decision Tree Confusion Matrix:  
[[55223 31121]  
[29768 56654]]

-----  
Training Random Forest...

|                                      |           |        |          |         |  |
|--------------------------------------|-----------|--------|----------|---------|--|
| Random Forest Classification Report: |           |        |          |         |  |
|                                      | precision | recall | f1-score | support |  |
| 0                                    | 0.71      | 0.78   | 0.74     | 86344   |  |
| 1                                    | 0.75      | 0.68   | 0.72     | 86422   |  |



|              |      |      |      |        |
|--------------|------|------|------|--------|
| accuracy     |      |      | 0.73 | 172766 |
| macro avg    | 0.73 | 0.73 | 0.73 | 172766 |
| weighted avg | 0.73 | 0.73 | 0.73 | 172766 |

Random Forest Confusion Matrix:

```
[[67032 19312]
 [27461 58961]]
```

-----  
Training Gradient Boosting (XGBoost)...

Gradient Boosting (XGBoost) Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.66      | 0.95   | 0.78     | 86344   |
| 1            | 0.91      | 0.50   | 0.65     | 86422   |
| accuracy     |           |        | 0.73     | 172766  |
| macro avg    | 0.78      | 0.73   | 0.71     | 172766  |
| weighted avg | 0.78      | 0.73   | 0.71     | 172766  |

Gradient Boosting (XGBoost) Confusion Matrix:

```
[[81970 4374]
 [42907 43515]]
```

-----  
Training Naive Bayes...

Naive Bayes Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.59      | 0.48   | 0.53     | 86344   |
| 1            | 0.56      | 0.67   | 0.61     | 86422   |
| accuracy     |           |        | 0.57     | 172766  |
| macro avg    | 0.58      | 0.57   | 0.57     | 172766  |
| weighted avg | 0.58      | 0.57   | 0.57     | 172766  |

Naive Bayes Confusion Matrix:

```
[[41453 44891]
 [28861 57561]]
```

-----  
Training LightGBM...

```
[LightGBM] [Info] Number of positive: 201521, number of negative: 201599
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.034312 seconds.
```

You can set `force\_row\_wise=true` to remove the overhead.

And if memory is not enough, you can set `force\_col\_wise=true`.

```
[LightGBM] [Info] Total Bins 10597
[LightGBM] [Info] Number of data points in the train set: 403120, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499903 -> initscore=-0.000387
[LightGBM] [Info] Start training from score -0.000387
```

LightGBM Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.66      | 0.94   | 0.77     | 86344   |
| 1            | 0.89      | 0.52   | 0.65     | 86422   |
| accuracy     |           |        | 0.73     | 172766  |
| macro avg    | 0.77      | 0.73   | 0.71     | 172766  |
| weighted avg | 0.77      | 0.73   | 0.71     | 172766  |

LightGBM Confusion Matrix:

```
[[80736 5608]
 [41644 44778]]
```

|         | Model                       | Accuracy | Precision | Recall   | F1-Score | \ |
|---------|-----------------------------|----------|-----------|----------|----------|---|
| 0       | Logistic Regression         | 0.571889 | 0.572065  | 0.571889 | 0.571646 |   |
| 1       | Decision Tree               | 0.647564 | 0.647597  | 0.647564 | 0.647541 |   |
| 2       | Random Forest               | 0.729270 | 0.731340  | 0.729270 | 0.728672 |   |
| 3       | Gradient Boosting (XGBoost) | 0.726329 | 0.782592  | 0.726329 | 0.712035 |   |
| 4       | Naive Bayes                 | 0.573110 | 0.575681  | 0.573110 | 0.569385 |   |
| 5       | LightGBM                    | 0.726497 | 0.774259  | 0.726497 | 0.714086 |   |
| AUC-ROC |                             |          |           |          |          |   |
| 0       |                             | 0.596138 |           |          |          |   |
| 1       |                             | 0.647555 |           |          |          |   |
| 2       |                             | 0.793471 |           |          |          |   |
| 3       |                             | 0.788423 |           |          |          |   |
| 4       |                             | 0.596682 |           |          |          |   |
| 5       |                             | 0.789205 |           |          |          |   |

# Hyper-parameter Tuning

Step 1

Step 2

Step 3

Step 4

Step 5.1

Step 5.2

Step 5.3

Step 5.1.1

Fitting 3 folds for each of 5 candidates, totalling 15 fits

[CV] END max\_depth=20, min\_samples\_leaf=4, min\_samples\_split=5, n\_estimators=300; total time= 1.8m in

[CV] END max\_depth=20, min\_samples\_leaf=4, min\_samples\_split=5, n\_estimators=300; total time= 1.8m in

[CV] END max\_depth=20, min\_samples\_leaf=4, min\_samples\_split=5, n\_estimators=300; total time= 1.8m in

[CV] END max\_depth=None, min\_samples\_leaf=2, min\_samples\_split=2, n\_estimators=200; total time= 8.9min

[CV] END max\_depth=None, min\_samples\_leaf=2, min\_samples\_split=2, n\_estimators=200; total time= 1.4min

[CV] END max\_depth=None, min\_samples\_leaf=2, min\_samples\_split=2, n\_estimators=200; total time= 1.4min

[CV] END max\_depth=None, min\_samples\_leaf=1, min\_samples\_split=5, n\_estimators=200; total time= 2.0min

[CV] END max\_depth=None, min\_samples\_leaf=1, min\_samples\_split=5, n\_estimators=200; total time= 1.4min

[CV] END max\_depth=None, min\_samples\_leaf=1, min\_samples\_split=5, n\_estimators=200; total time= 1.4min

[CV] END max\_depth=30, min\_samples\_leaf=1, min\_samples\_split=2, n\_estimators=300; total time=44.1m in

[CV] END max\_depth=30, min\_samples\_leaf=1, min\_samples\_split=2, n\_estimators=300; total time= 2.6m in

[CV] END max\_depth=30, min\_samples\_leaf=1, min\_samples\_split=2, n\_estimators=300; total time= 2.5m in

[CV] END max\_depth=20, min\_samples\_leaf=1, min\_samples\_split=10, n\_estimators=300; total time= 1.9

```
min
[CV] END max_depth=20, min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 2.3
min
[CV] END max_depth=20, min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 2.3
min
Step 5.2.1

Fitting 3 folds for each of 5 candidates, totalling 15 fits

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:11] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.6; total time= 4.4s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:14] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.6; total time= 2.4s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:16] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.6; total time= 5.8s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:22] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=3, n_estimators=100, subsample=0.6; total time= 2.7s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:25] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=3, n_estimators=100, subsample=0.6; total time= 2.4s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:28] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=3, n_estimators=100, subsample=0.6; total time= 2.9s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:31] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=200, subsample=0.8; total time= 11.5s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:42] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

    warnings.warn(smsg, UserWarning)
```

```

[CV] END learning_rate=0.01, max_depth=10, n_estimators=200, subsample=0.8; total time= 9.2s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:10:52] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=200, subsample=0.8; total time= 7.9s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:11:00] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.8; total time= 3.4s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:11:03] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.8; total time= 3.0s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:11:06] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.8; total time= 2.9s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:11:09] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=100, subsample=0.6; total time= 4.3s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:11:14] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=100, subsample=0.6; total time= 4.4s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:11:18] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=100, subsample=0.6; total time= 4.2s
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [10:11:23] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)

```

#### Step 5.3.1

```

Fitting 3 folds for each of 5 candidates, totalling 15 fits
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.022893 sec
onds.
You can set `force_row_wise=true` to remove the overhead.

```

And if memory is not enough, you can set `force\_col\_wise=true`.

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=gbdt, learning_rate=0.2, n_estimators=100, num_leaves=50; total time= 1.9
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010730 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=gbdt, learning_rate=0.2, n_estimators=100, num_leaves=50; total time= 1.9
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.011540 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=gbdt, learning_rate=0.2, n_estimators=100, num_leaves=50; total time= 1.6
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008373 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=50; total time= 11.8
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010952 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=50; total time= 11.7
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.012068 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=50; total time= 11.9
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008941 seconds.
You can set `force_row_wise=true` to remove the overhead.
```

And if memory is not enough, you can set `force\_col\_wise=true`.

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=31; total time= 9.3
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010446 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=31; total time= 10.3
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.011120 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=31; total time= 9.9
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.007598 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=gbd, learning_rate=0.1, n_estimators=200, num_leaves=31; total time= 2.6
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.026696 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=gbd, learning_rate=0.1, n_estimators=200, num_leaves=31; total time= 3.6
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.011196 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=gbd, learning_rate=0.1, n_estimators=200, num_leaves=31; total time= 2.6
s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010616 seconds.
You can set `force_row_wise=true` to remove the overhead.
```

```

And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.1, n_estimators=300, num_leaves=100; total time= 26.6s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.013189 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.1, n_estimators=300, num_leaves=100; total time= 23.4s
[LightGBM] [Info] Number of positive: 67028, number of negative: 153648
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010147 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 220676, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
[CV] END boosting_type=dart, learning_rate=0.1, n_estimators=300, num_leaves=100; total time= 27.0s
[LightGBM] [Info] Number of positive: 100542, number of negative: 230472
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.024448 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 347
[LightGBM] [Info] Number of data points in the train set: 331014, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.303739 -> initscore=-0.829554
[LightGBM] [Info] Start training from score -0.829554
Best model: Random Forest with F1 score: 0.5046067373878573
RandomForestClassifier(max_depth=30, n_estimators=300, random_state=42)

```

## Best Model

```
RandomForestClassifier(max_depth=30, n_estimators=300, random_state=42)
```

Step 1

Step 2

Step 3

Step 4

Step 5.1

Step 5.2

Step 5.3

Step 5.1.1

Fitting 3 folds for each of 5 candidates, totalling 15 fits

```
[CV] END max_depth=20, min_samples_leaf=4, min_samples_split=5, n_estimators=300; total time= 3.9min
```

```
[CV] END max_depth=20, min_samples_leaf=4, min_samples_split=5, n_estimators=300; total time= 4.2min
```

```
[CV] END max_depth=20, min_samples_leaf=4, min_samples_split=5, n_estimators=300; total time= 4.3min
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=200; total time= 3.2min
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=200; total time= 3.4min
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=200; total time= 3.4min
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=5, n_estimators=200; total time= 3.9min
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=5, n_estimators=200; total time= 3.3min
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=5, n_estimators=200; total time= 3.0min
[CV] END max_depth=30, min_samples_leaf=1, min_samples_split=2, n_estimators=300; total time= 5.2min
[CV] END max_depth=30, min_samples_leaf=1, min_samples_split=2, n_estimators=300; total time= 4.4min
[CV] END max_depth=30, min_samples_leaf=1, min_samples_split=2, n_estimators=300; total time= 4.1min
[CV] END max_depth=20, min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 3.3min
[CV] END max_depth=20, min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 3.3min
[CV] END max_depth=20, min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 3.4min
```

#### Step 5.2.1

Fitting 3 folds for each of 5 candidates, totalling 15 fits

```
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:40:26] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

```
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.6; total time= 5.1s
```

```
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:40:31] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

```
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.6; total time= 4.8s
```

```
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:40:36] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

```
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.6; total time= 6.3s
```

```
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:40:42] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

```
[CV] END learning_rate=0.01, max_depth=3, n_estimators=100, subsample=0.6; total time= 4.2s
```

```
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:40:46] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

```
[CV] END learning_rate=0.01, max_depth=3, n_estimators=100, subsample=0.6; total time= 3.6s
```



```
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:40:50] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=3, n_estimators=100, subsample=0.6; total time= 4.1s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:40:55] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=200, subsample=0.8; total time= 27.5s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:41:22] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=200, subsample=0.8; total time= 21.1s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:41:43] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=200, subsample=0.8; total time= 18.5s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:42:02] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.8; total time= 5.7s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:42:08] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.8; total time= 7.3s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:42:15] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.1, max_depth=3, n_estimators=200, subsample=0.8; total time= 6.0s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:42:21] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=100, subsample=0.6; total time= 9.8s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:42:31] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=100, subsample=0.6; total time= 9.9s
```

```
C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:42:41] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
[CV] END learning_rate=0.01, max_depth=10, n_estimators=100, subsample=0.6; total time= 10.0s

C:\Users\admin\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [11:42:51] WARNING:
C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0015a694724fa8361-1\xgboost\xg
boost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

### Step 5.3.1

```
Fitting 3 folds for each of 5 candidates, totalling 15 fits
[LightGBM] [Info] Number of positive: 153463, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.047518 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10871
[LightGBM] [Info] Number of data points in the train set: 307138, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499655 -> initscore=-0.001380
[LightGBM] [Info] Start training from score -0.001380
[CV] END boosting_type=gbd, learning_rate=0.2, n_estimators=100, num_leaves=50; total time= 3.1
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153676
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.049018 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10867
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499653 -> initscore=-0.001387
[LightGBM] [Info] Start training from score -0.001387
[CV] END boosting_type=gbd, learning_rate=0.2, n_estimators=100, num_leaves=50; total time= 3.1
s
[LightGBM] [Info] Number of positive: 153464, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.040361 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10863
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499657 -> initscore=-0.001374
[LightGBM] [Info] Start training from score -0.001374
[CV] END boosting_type=gbd, learning_rate=0.2, n_estimators=100, num_leaves=50; total time= 3.0
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.048754 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10871
[LightGBM] [Info] Number of data points in the train set: 307138, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499655 -> initscore=-0.001380
[LightGBM] [Info] Start training from score -0.001380
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=50; total time= 27.8
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153676
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.045380 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
```

```

[LightGBM] [Info] Total Bins 10867
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499653 -> initscore=-0.001387
[LightGBM] [Info] Start training from score -0.001387
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=50; total time= 27.6
s
[LightGBM] [Info] Number of positive: 153464, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.059656 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10863
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499657 -> initscore=-0.001374
[LightGBM] [Info] Start training from score -0.001374
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=50; total time= 28.1
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.047733 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10871
[LightGBM] [Info] Number of data points in the train set: 307138, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499655 -> initscore=-0.001380
[LightGBM] [Info] Start training from score -0.001380
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=31; total time= 26.9
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153676
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.052242 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10867
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499653 -> initscore=-0.001387
[LightGBM] [Info] Start training from score -0.001387
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=31; total time= 25.1
s
[LightGBM] [Info] Number of positive: 153464, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.052547 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10863
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499657 -> initscore=-0.001374
[LightGBM] [Info] Start training from score -0.001374
[CV] END boosting_type=dart, learning_rate=0.2, n_estimators=200, num_leaves=31; total time= 38.0
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.057067 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10871
[LightGBM] [Info] Number of data points in the train set: 307138, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499655 -> initscore=-0.001380
[LightGBM] [Info] Start training from score -0.001380
[CV] END boosting_type=gbd, learning_rate=0.1, n_estimators=200, num_leaves=31; total time= 7.8
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153676
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.057011 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

```

```

[LightGBM] [Info] Total Bins 10867
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499653 -> initscore=-0.001387
[LightGBM] [Info] Start training from score -0.001387
[CV] END boosting_type=gbdt, learning_rate=0.1, n_estimators=200, num_leaves=31; total time= 7.7
s
[LightGBM] [Info] Number of positive: 153464, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.054723 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10863
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499657 -> initscore=-0.001374
[LightGBM] [Info] Start training from score -0.001374
[CV] END boosting_type=gbdt, learning_rate=0.1, n_estimators=200, num_leaves=31; total time= 7.4
s
[LightGBM] [Info] Number of positive: 153463, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.062695 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10871
[LightGBM] [Info] Number of data points in the train set: 307138, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499655 -> initscore=-0.001380
[LightGBM] [Info] Start training from score -0.001380
[CV] END boosting_type=dart, learning_rate=0.1, n_estimators=300, num_leaves=100; total time= 1.5m
in
[LightGBM] [Info] Number of positive: 153463, number of negative: 153676
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.058857 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10867
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499653 -> initscore=-0.001387
[LightGBM] [Info] Start training from score -0.001387
[CV] END boosting_type=dart, learning_rate=0.1, n_estimators=300, num_leaves=100; total time= 1.3m
in
[LightGBM] [Info] Number of positive: 153464, number of negative: 153675
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.047332 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10863
[LightGBM] [Info] Number of data points in the train set: 307139, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499657 -> initscore=-0.001374
[LightGBM] [Info] Start training from score -0.001374
[CV] END boosting_type=dart, learning_rate=0.1, n_estimators=300, num_leaves=100; total time= 1.1m
in
[LightGBM] [Info] Number of positive: 230195, number of negative: 230513
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.063700 sec
onds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 10867
[LightGBM] [Info] Number of data points in the train set: 460708, number of used features: 45
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499655 -> initscore=-0.001380
[LightGBM] [Info] Start training from score -0.001380
Best model: Random Forest with F1 score: 0.7489743037747075
RandomForestClassifier(min_samples_split=5, n_estimators=200, random_state=42)

```

## Random Forest Classifier gives best performance with F1

**score: 0.74**

We can export this model in .pkl file to be used for prediction